

2027 UF Guide to AI in Science

Nicholas A. Coles Abhay Alaukik Majed A. Arishi
Shanting Chen Annabel V. Dang Jungjae Im Nabiha Iqbal
Minyi Jiang Rui Jin Zhongchi Li Sooyeon Park
Mariana G. Rebello Jihyun Soh Paola J. Sullivan
Lyra G. Toohey Sofia C. Villanueva

2026-09-16

Table of contents

Ideation and Project Management	6
Ideation and Project Management: Using AI	7
Areas of consensus — and dissensus	8
Ideation and Project Management: Critiquing AI	10
Areas of consensus — and dissensus	10
 Literature Review and Synthesis	 13
Literature Review and Synthesis: Using AI	14
Areas of consensus — and dissensus	14
Literature Review and Synthesis: Critiquing AI	17
 Data Collection and Generation	 18
Data Collection and Generation: Using AI	19
Data Collection and Generation: Critiquing AI	20
 Analysis and Modeling	 21
Analysis and Modeling: Using AI	22
Analysis and Modeling: Critiquing AI	23
 Writing and Communication	 24
Writing and Communication: Using AI	25
Writing and Communication: Critiquing AI	26

Ethics and Outreach	27
Ethics and Outreach: Using AI	28
Ethics and Outreach: Critiquing AI	29
Personal AI Philosophy Statements	30
References	31

This was written by humans.

But we provide links to helpful AI-generated content along the way.

AI is radically redefining the nature of scientific inquiry — for better and for worse.

In Fall 2026, faculty, graduate students, and postdocs across multiple departments at the University of Florida came together to understand this urgent development. In a seminar titled *Collective [Artificial] Intelligence in Science*, we documented and discussed how AI is changing each component of our research workflow. This includes project management, idea development, data analysis, writing, and dissemination. To push our understanding, we alternated between the perspectives of an AI zealot and an AI critic — identifying key resources, consensus statements, and areas of disagreement along the way. The result is the **2027 UF Guide to AI in Science**.

The seminar began with an introduction to big-team science — its rise¹, competitive advantages², and relevance to the current moment. For instance, the paper describing OpenAI’s GPT-1 had 4 authors³; the paper describing GPT-3 had 31 authors⁴; and the paper describing GPT-5 had 483 authors⁵ (Figure 1). In the artificial realm, a prominent AI-led cybersecurity attack on the online community, Hugging Face, involved coordinated action among not one but *thousands* of AI agents⁶. Similarly, over 10,000 AI agents recently exchanged approximately 5 million messages over 88 hours in order to solve the Navier–Stokes equations⁷ — one of the seven Clay Millennium Prize Problems. Whether it be human or artificial, some of our most impressive accomplishments are requiring something rarely formally taught in science: large-scale collaboration⁸.

After reviewing this backdrop, the course itself became a large-scale collaboration, dedicated to understanding how AI is changing the nature of science. Each week followed a consistent structure (Figure 2). First, we would select a component of the research pipeline to examine (e.g., data analysis). Second, we would spend one week compiling resources for understanding the emerging role of AI in that part of the research pipeline. Third, we would engage in peer-to-peer teaching. Last, we would use a live online voting platform to solicit, debate, and vote upon key resources and takeaways for the week^{9,10}. Through these exercises, seminar participants received hands on experience with large-scale collaboration¹¹. Along the way, not only developed and debated their *personal* philosophy statements – but also the general




Figure 1: Credited co-authors on flagship AI model reports from OpenAI, Google, and Meta, 2018–2025. One square is one credited person; each block’s area is its author list.

direction of AI at UF. Their results culminate into this **2027 UF Guide to AI in Science**: a summary of our top resources, takeaways, and disagreements.




Figure 2: Structure of the UF Fall 2026 seminar, titled Collective [Artificial] Intelligence in Science.

Ideation and Project Management


Ideation and Project Management: Using AI

This was written by humans.

But we also like this [podcast](#) (20:36) of the entire 2-week unit on Ideation and Project Management, generated with Gemini Notebook¹².

Many researchers' first contact with AI tools is in the realm of ideation or project management (Figure 1). For example, an in-class demo focused on an AI-assisted workflow for managing research trainee projects. More specifically, one seminar participant has their trainees submit weekly updates to a folder shared with the LLM, Claude¹³ (Figure 2, Panel A). This allows them to later prompt the model with questions about past conversations, challenges, decisions, and solutions (Figure 2, Panel B). Specifically, the seminar participant demonstrated how Claude can be used to efficiently take stock of a trainee's progress, have a conversation with the LLM about potential future directions, and generate a simulated figure to send back to the trainee to convey next steps (Figure 2, Panel C).

Top discussion points



Top resources




Figure 1: Overview of Week 2 of the Fall 2026 Collective [Artificial] Intelligence Seminar. Left: top discussion points, along with estimated level of agreement among seminar participants. Right: three most popular resources identified by seminar participants.

The resources compiled by seminar participants illustrate that AI is a tool that is increasingly being used to expand ideation and project management capabilities (Figure 1). For example, as early as 2025, researchers documented that LLMs were capable of generating research ideas that were judged more novel than those of human experts¹⁴. These capabilities can be further improved by integrating causal discovery algorithms. More specifically, LLMs can efficiently extract which causal relationships have been reported in past research. When mapped across studies, causal discovery algorithms can be used to identify gaps in causal network maps — i.e., potential future research directions¹⁵. Such results highlight the increasing sophistication with which researchers are deploying AI to organize information, extract ideas, and even develop theory¹⁶.

Areas of consensus — and dissensus

Disagreement is common when researchers collaborate^{9,10} — and our seminar discussion was marked by both consensus and dissensus.

In terms of consensus, there was broad agreement that LLMs can assist with ideation, project management, and the organization of information. There was also broad agreement that this is changing the nature of competition in science — placing more emphasis on who can better leverage these tools for ideation and execution.

Yet, there was clear disagreement about how far things should go. Seminar participants strongly disagreed about whether AI is reliably superior at ideation, whether the institution of science should shift to more AI-led work, and whether humanity should strive for fully autonomous AI scientists. These disagreements echo throughout the scientific landscape, with recent back-to-back *Nature* pieces on both the capabilities¹⁷ and limitations¹⁸ of science with fewer humans and more AI. These limitations became the focus of the following week of the Fall 2026 Collective [Artificial] Intelligence Seminar (see Section).

The screenshot shows a Jupyter Notebook interface with a dark theme. The top bar indicates the file is named '04_21_2020.ipynb' and is located in a directory named '04_21_2020.ipynb - chadon'. The notebook is open to a cell titled 'Algorithm for Modal Country'. The code in the cell is as follows:

```

import pandas as pd
import numpy as np
import sklearn

# Load the data
data = pd.read_csv('data/04_21_2020.csv')

# Preprocess the data
data = data.dropna()
data = data.reset_index()

# Split the data into training and testing sets
train_data, test_data = sklearn.model_selection.train_test_split(data, random_state=42)

# Train the model
model = sklearn.ensemble.RandomForestClassifier()
model.fit(train_data[['Year', 'Country']], train_data['Mode'])

# Test the model
test_data[['Year', 'Country']].head()
model.predict(test_data[['Year', 'Country']]).head()

```

The notebook interface includes a file explorer on the left showing the file structure, a top bar with navigation and search options, and a bottom status bar showing the current cell's execution time and memory usage.

```
Claude Code v2.1.245
Opus 5 · Claude Max
~/Documents/GitHub/LabManual/checkin/Cindy

> /model
  | Set model to Opus 5 (1M context) and saved as your default for new sessions

> run through cindy's notes and give me a sense for how she is identifying "DEI scholars"

• I'll look through Cindy's check-in notes.

  Listed 1 directory, ran 5 shell commands (ctrl+o to expand)

• Read all 19 of Cindy's check-ins (04/01-08/18). Here's how "DEI scholar" actually gets built, and where it wobbles.

  It's a two-stage construct: paper label → author score → cutoff
```

Texas

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 4.2 \times 10^3$, $BF_{1,2} = 8.1 \times 10^3$, $BF_{1,3} = 2.6 \times 10^3$.

Florida

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 1.8 \times 10^3$, $BF_{1,2} = 3.4 \times 10^3$, $BF_{1,3} = 170$.

Tennessee

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 29$, $BF_{1,2} = 4.1$, $BF_{1,3} = 1.2$.

United States

Line graph: % outmigrant (2018-2025). BF values: $BF_{1,0} = 140$, $BF_{1,2} = 0.3$, $BF_{1,3} = 0.42$.

Legend: Texas (orange), Florida (green), Tennessee (pink), New York (dark blue), California (yellow), United States (dark blue), Canada (teal), United Kingdom (maroon).

9

Ideation and Project Management: Critiquing AI

This was written by humans.

But we also like this [podcast](#) (20:36) of the entire 2-week unit on Ideation and Project Management, generated with Gemini Notebook¹².

In February 2026, an engineer running an online educational program used Claude Code¹⁹ to help launch a new side project. Instead of launching something new, they ultimately deleted everything they had already built: servers, networks, and years' worth of datasets. The failure point, the engineer would later suggest, was over-automation²⁰. The engineer never checked whether backups could be restored; they never pushed back on questionable recommendations from the coding agent; and they allowed it to run irreversible commands completely unsupervised. Issues like these are bound to become more common as scientists increasingly rely on AI — the focus of Week 3 of the Fall 2026 Collective [Artificial] Intelligence Seminar (Figure 1).


This week's most popular resources focused on the pitfalls of using AI for ideation and project management (Figure 1). For example, one study used electroencephalography to measure brain connectivity patterns while writing alone, with a search engine, or with an LLM. Results suggested that LLM users exhibited the lowest levels of distributed brain activity — an index of lower cognitive engagement²¹. Similarly, when researchers rely more on AI for routine tasks like reproducibility checks, they catch fewer errors, propose weaker robustness checks, and ultimately move slower²². As scientists continue this shift towards AI, collective implications emerge. For instance, an analysis of 41 million research papers suggests that the adoption of AI tools and methods increases productivity, impact, and career outcomes for individuals — but at the expense of an overall narrowing of the scope of science²³ (Figure 2).

Areas of consensus — and dissensus

Particularly after focusing a prior week on the benefits of AI, seminar discussion was marked by areas of consensus and dissensus^{9,10}.

In general, there was strong agreement that concerns about AI have evolved beyond hallucination proneness²⁴. Decreased vigilance and a collective narrowing of science received more attention

Top discussion points



Top resources



Your Brain on ChatGPT

54 essay writers under EEG: LLM users showed the weakest neural connectivity and least essay ownership.

Identified by Majed A. Arishi and Rui Jin



AI Tools Expand Impact, Contract Focus

41.3 million papers: AI-augmented scientists publish and are cited far more, while topics studied narrow.

Identified by Zhongchi Li and Sooyeon Park



AI-Assisted vs. AI-Led Research Teams

Human-only and AI-assisted teams reproduced over 90% of published results; AI-led teams managed 37%.

Identified by Paola J. Sullivan and Mariana G. Rebello

Figure 1: Overview of Week 3 of the Fall 2026 Collective [Artificial] Intelligence Seminar. Left: top discussion points, along with estimated level of agreement among seminar participants. Right: three most popular resources identified by seminar participants.

among seminar participants — with near-unanimous agreement that maximizing value from AI requires training, practice, and refinement²⁵. Yet, there was clear disagreement about the magnitude of the issue. Does AI hinder more than help with critical thinking? In the midst of widespread concerns about the robustness of existing research practices²⁶, to what extent should we aim for increases in productivity? Last, is a narrowing of the scope of science necessarily a bad tradeoff²⁷? After all, depression, addiction, cancer, dementia, climate change, and polarization are arguably enduring problems in need of sustained focus. Whether those problems are solved more quickly through increases in individual productivity, decreases in collective scope, and a faster overall science remains an open question.




Figure 2: Illustration of the narrowing of the scope of science, created by Zhongchi Li, Sooyeon Park, and Nicholas Coles, using Claude¹³, Gemini Notebook, and ChatGPT²⁸.

Literature Review and Synthesis

Literature Review and Synthesis: Using AI

This content was identified, synthesized, and written by humans.

But we also like this AI-generated review (TBD) of the benefits and drawbacks of AI for literature review and synthesis.


LLMs have an unprecedented capacity to conduct and assist with literature reviews — making them arguably one of the most essential additions to the modern-day research workflow (Figure 1). This week, an in-class demo focused on how the “days of research assistants, Google Alerts, and journal subscriptions” are being replaced by AI-assisted workflows (Figure 2). For example, armed with markdown files for organizing search strategies, key insights, and overarching project goals, one demo deployed Claude¹³ to (a) search, screen, and organize resources, (b) extract key insights, connections, and themes, (c) brief the PI on recent developments, and (d) adjust the project vision based on PI feedback (Figure 2, Panel B). This process runs automatically each week, allowing this researcher to sort through an increasingly growing and unmanageable scientific literature with more efficiency than ever before.

From custom workflows (Figure 2) to LLM skills²⁹ to off-the-shelf tools like Consensus³⁰, resources compiled by seminar participants painted an overwhelming picture: AI *already* has revolutionized the literature review space. For example, an early evaluation demonstrated that LLMs are capable of reviewing up to 110 research records per minute, often achieving accuracy comparable to human coders who need months for the same task³¹. However, not all models are created equal — some performing at human-level accuracy and others landing at the floor³². This has led to new programs of research dedicated to optimizing these AI systems. For example, AI agents are capable of performing more accurate literature reviews when allowed to work in teams³³ — an insight that parallels many of the collective intelligence themes reviewed in this seminar^{1,2,8} (see [Preface](#)).

Areas of consensus — and dissensus

Working in a team setting, of course, requires navigating both consensus and dissensus^{9,10}. In terms of consensus, there was strong agreement that AI tools are ubiquitous in literature review, with browser-based applications like Consensus³⁰, Gemini Notebook¹², and ChatGPT²⁸ providing non-technical entry points for nearly all researchers. Furthermore, there was agreement that such workflows are still relatively novel, indicating a strong need for continued evaluation

Top discussion points



Top resources



Full Agreement Across Multiple LLMs

ChatGPT/Gemini/Claude all had to agree before a screening call stood – cutting manual screening 70%.

Identified by Nabihah Iqbal and Minyi Jiang



How Well Do LLMs Screen Abstracts?

GPT-4 matched human include/exclude calls 84–95% of the time and cleared 200 abstracts in 20 minutes.

Identified by Jihyun Soh and Jungjae Im



Human-Comparable Screening Sensitivity

Splitting eligibility into three questions – design, population, controls – brought GPT to human sensitivity.

Identified by Annabel V. Dang and Lyra G. Toohey

Figure 1: Overview of Week 4 of the Fall 2026 Collective [Artificial] Intelligence Seminar. Left: top discussion points, along with estimated level of agreement among seminar participants. Right: three most popular resources identified by seminar participants.

and reflection. As part of this reflection, we reached our most poignant disagreements. Are systematic reviews still noteworthy contributions to science? Are we ready to rely upon these tools in our own research? And, if so, what’s left for humans?

“What’s left for humans?” was not an abstract question. Nearly every seminar participant considered literature reviews an essential part of their initial training. Yet, many also acknowledged that they’ve largely stopped offering these same training opportunities in their own labs. In the ultra-competitive and fast-moving world of science³⁴, many researchers simply feel they cannot afford to lose the speed advantage of LLM-assisted literature review. Seminar participants had mixed emotions³⁵ about these developments. For example, some seminar participants noted the extremely unpleasant nature of a literature review assignment. One participant likened the notion of having a human spend months on a literature review task an LLM can complete in hours to “academic hazing”. Some suggested it’s a necessary skill to develop, but not to personally practice – akin to learning long division before being handed a calculator. Yet, even if that is the case, a philosophical tension remained. Every academic in the seminar has been collectively entrusted by society to understand, maintain, and develop the scientific record. What does it mean when we are increasingly disengaging from that literature directly? To what extent had we already done so far before the arrival of AI? Themes such as this became the focus of the following week of the Fall 2026 Collective [Artificial] Intelligence Seminar (see Section).




Figure 2: A seminar participant shared an AI-assisted workflow for literature review and synthesis (Panel A) and demonstrated it running against a working project (Panel B). The agent searches and screens against a written strategy, retrieves and organizes the papers, extracts insights, connections, and themes into separate files, and then talks the results through with the principal investigator - who updates the project’s guiding documents before the loop repeats.

Literature Review and Synthesis: Critiquing AI

This content was identified, synthesized, and written by humans.

But we also like this AI-generated review (TBD) of the benefits and drawbacks of AI for literature review and synthesis.

Covered in Week 5 of the seminar (September 17, 2026). Estimated date of arrival: September 27, 2026.

Data Collection and Generation

Data Collection and Generation: Using AI

Covered in Week 6 of the seminar (September 24, 2026). Estimated date of arrival: October 4, 2026.

Data Collection and Generation: Critiquing AI

Covered in Week 7 of the seminar (October 1, 2026). Estimated date of arrival: October 11, 2026.

Analysis and Modeling

Analysis and Modeling: Using AI

Covered in Week 8 of the seminar (October 8, 2026). Estimated date of arrival: October 18, 2026.

Analysis and Modeling: Critiquing AI

Covered in Week 9 of the seminar (October 15, 2026). Estimated date of arrival: October 25, 2026.

Writing and Communication

Writing and Communication: Using AI

Covered in Week 10 of the seminar (October 22, 2026). Estimated date of arrival: November 1, 2026.

Writing and Communication: Critiquing AI

Covered in Week 11 of the seminar (October 29, 2026). Estimated date of arrival: November 8, 2026.

Ethics and Outreach

Ethics and Outreach: Using AI

Covered in Week 12 of the seminar (November 5, 2026). Estimated date of arrival: November 15, 2026.

Ethics and Outreach: Critiquing AI

Covered in Week 13 of the seminar (November 12, 2026). Estimated date of arrival: November 22, 2026.

Personal AI Philosophy Statements

Covered in Week 14 of the seminar (November 19, 2026). Estimated date of arrival: November 29, 2026.

References

1. Coles, N. A., Hamlin, J. K., Sullivan, L. L., Parker, T. H. & Altschul, D. [Build up big-team science](#). *Nature* **601**, 505–507 (2022).
2. Coles, N. A., Goes Braga Takayanagi, J. F., Grant, G. L. & Basnight-Brown, D. M. [The rise, impact, and imbalances of big-team psychology](#). *Psychological Science* **37**, 287–297 (2026).
3. Radford, A., Narasimhan, K., Salimans, T. & Sutskever, I. *Improving Language Understanding by Generative Pre-Training*. https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf (2018).
4. Brown, T. B., Mann, B., Ryder, N., Subbiah, M., *et al.* Language models are few-shot learners. *arXiv preprint* <https://doi.org/10.48550/arXiv.2005.14165> (2020) doi:10.48550/arXiv.2005.14165.
5. OpenAI. OpenAI GPT-5 system card. *arXiv preprint* <https://doi.org/10.48550/arXiv.2601.03267> (2025) doi:10.48550/arXiv.2601.03267.
6. OpenAI. The hugging face incident and the road ahead. <https://openai.com/index/hugging-face-incident-and-the-road-ahead/> (2026).
7. OpenAI. On the Navier–Stokes Millennium Prize Problem. <https://openai.com/index/navier-stokes-solution/> (2026).
8. National Research Council. *Enhancing the Effectiveness of Team Science*. (National Academies Press, Washington, DC, 2015). doi:10.17226/19007.
9. Coles, N. A. *et al.* [Team scientists should normalize disagreement](#). *Science* **384**, 1076–1077 (2024).
10. Dalkey, N. & Helmer, O. [An experimental application of the Delphi method to the use of experts](#). *Management Science* **9**, 458–467 (1963).
11. Baumgartner, H. A. *et al.* [How to build up big team science: A practical guide for large-scale collaborations](#). *Royal Society Open Science* **10**, 230235 (2023).
12. Google. Gemini notebook. <https://notebook.google.com> (2026).
13. Anthropic. Claude. <https://claude.ai> (2026).
14. Si, C., Yang, D. & Hashimoto, T. [Can LLMs generate novel research ideas? A large-scale human study with 100+ NLP researchers](#). in *International conference on learning representations* (eds Yue, Y., Garg, A., Peng, N., Sha, F. & Yu, R.) vol. 2025 94003–94092 (2025).
15. Tong, S., Mao, K., Huang, Z., Zhao, Y. & Peng, K. Automating psychological hypothesis generation with AI: When large language models meet causal graph. *Humanities and Social Sciences Communications* **11**, 896 (2024).

16. Schmitt, B., Hao, S., Pham, M. T. & Hofstetter, R. Generative grounded theory (GGT): Inductive theory building in the age of generative AI. *Journal of Consumer Psychology* <https://doi.org/10.1002/jcpy.70036> (2026) doi:10.1002/jcpy.70036.
17. Yamada, Y. *et al.* Towards end-to-end automation of AI research. *Nature* **651**, 914–919 (2026).
18. Nature. Why AI cannot do good science without humans. *Nature* <https://www.nature.com/articles/d41586-026-01551-3> (2026).
19. Anthropic. Claude code. <https://claude.com/product/claude-code> (2026).
20. Grigorev, A. Claude Code wiped our production database with a Terraform command. https://x.com/Al_Grigor/status/2029889772181934425 (2026).
21. Kosmyna, N. *et al.* Your brain on ChatGPT: Accumulation of cognitive debt when using an AI assistant for essay writing task. *arXiv preprint arXiv:2506.08872* <https://doi.org/10.48550/arXiv.2506.08872> (2025) doi:10.48550/arXiv.2506.08872.
22. Brodeur, A. *et al.* AI-assisted teams outperform AI-led teams but not human-only teams in assessing research reproducibility in quantitative social science. *Proceedings of the National Academy of Sciences* **123**, e2524747123 (2026).
23. Hao, Q., Xu, F., Li, Y. & Evans, J. Artificial intelligence tools expand scientists’ impact but contract science’s focus. *Nature* **649**, 1237–1243 (2026).
24. Topaz, M., Roguin, N., Gupta, P., Zhang, Z. & Peltonen, L.-M. Fabricated citations: An audit across 2 · 5 million biomedical papers. *The Lancet* **407**, 1779–1781 (2026).
25. Southworth, J. *et al.* Developing a model for AI across the curriculum: Transforming the higher education landscape via innovation in AI literacy. *Computers and Education: Artificial Intelligence* **4**, 100127 (2023).
26. Munafò, M. R. *et al.* A manifesto for reproducible science. *Nature Human Behaviour* **1**, 0021 (2017).
27. Kuhn, T. S. The essential tension: Tradition and innovation in scientific research. in *The essential tension: Selected studies in scientific tradition and change* 225–239 (University of Chicago Press, Chicago, 1977). doi:10.7208/9780226217239-010.
28. OpenAI. ChatGPT. <https://chatgpt.com> (2026).
29. Wu, E. C.-I. Academic research skills for Claude Code. <https://github.com/Imbad0202/academic-research-skills> (2026).
30. Consensus. Consensus. <https://consensus.app> (2026).
31. Matsui, K. *et al.* Human-comparable sensitivity of large language models in identifying eligible studies through title and abstract screening: 3-layer strategy using GPT-3.5 and GPT-4 for systematic reviews. *Journal of Medical Internet Research* **26**, e52758 (2024).
32. Li, M., Sun, J. & Tan, X. Evaluating the effectiveness of large language models in abstract screening: A comparative analysis. *Systematic Reviews* **13**, 219 (2024).
33. Hilkenmeier, F., Stoltenberg, M. & Stierle, C. Using full agreement across multiple large language models for title-and-abstract screening in systematic reviews: A proof-of-concept. *Systematic Reviews* **15**, 191 (2026).

34. Bornmann, L., Haunschild, R. & Mutz, R. [Growth rates of modern science: A latent piecewise growth curve approach to model publication numbers from established and new literature databases](#). *Humanities and Social Sciences Communications* **8**, 224 (2021).
35. Larsen, J. T., Coles, N. A. & Jordan, D. K. [Varieties of mixed emotional experience](#). *Current Opinion in Behavioral Sciences* **15**, 72–76 (2017).